

# plant-prediction

November 1, 2024

```
[61]: import numpy as np
import pandas as pd

import os
#for dirname, _, filenames in os.walk('C:/Users/KIIT0001/Desktop/archive/
↳medicinal/Indian Medicinal Leaves Image Datasets/medicinal_plant_dataset'):
#    for filename in filenames:
#        print(os.path.join(dirname, filename))
```

```
[6]: import tensorflow as tf
```

```
[7]: dataset = tf.keras.utils.image_dataset_from_directory("C:/Users/KIIT0001/
↳Desktop/archive/medicinal/Indian Medicinal Leaves Image Datasets/
↳medicinal_plant_dataset")
```

Found 5945 files belonging to 40 classes.

```
[ ]: ds_train = tf.keras.utils.image_dataset_from_directory(
    'C:/Users/KIIT0001/Desktop/archive/medicinal/Indian Medicinal Leaves Image_
↳Datasets/medicinal_plant_dataset',
    image_size=(224, 224),
    batch_size=16,
    validation_split=0.4,
    subset="training",
    seed=42
)
```

Found 5945 files belonging to 40 classes.

Using 3567 files for training.

```
[ ]: ds_val = tf.keras.utils.image_dataset_from_directory(
    'C:/Users/KIIT0001/Desktop/archive/medicinal/Indian Medicinal Leaves Image_
↳Datasets/medicinal_plant_dataset',
    image_size=(224, 224),
    batch_size=16,
    validation_split=0.4,
    subset="validation",
    seed=42
)
```

```
)
```

Found 5945 files belonging to 40 classes.  
Using 2378 files for validation.

```
[10]: import tensorflow_datasets as tfds
```

```
batch_size = 64
dataset_name = dataset
class_names = dataset.class_names

print(class_names)
```

```
['Aloevera', 'Amla', 'Amruta_Balli', 'Arali', 'Ashoka', 'Ashwagandha',
'Avacado', 'Bamboo', 'Basale', 'Betel', 'Betel_Nut', 'Brahmi', 'Castor',
'Curry_Leaf', 'Doddapatre', 'Ekka', 'Ganike', 'Gauva', 'Geranium', 'Henna',
'Hibiscus', 'Honge', 'Insulin', 'Jasmine', 'Lemon', 'Lemon_grass', 'Mango',
'Mint', 'Nagadali', 'Neem', 'Nithyapushpa', 'Nooni', 'Pappaya', 'Pepper',
'Pomegranate', 'Raktachandini', 'Rose', 'Sapota', 'Tulasi', 'Wood_sorel']
```

```
[11]: len(class_names)
```

```
[11]: 40
```

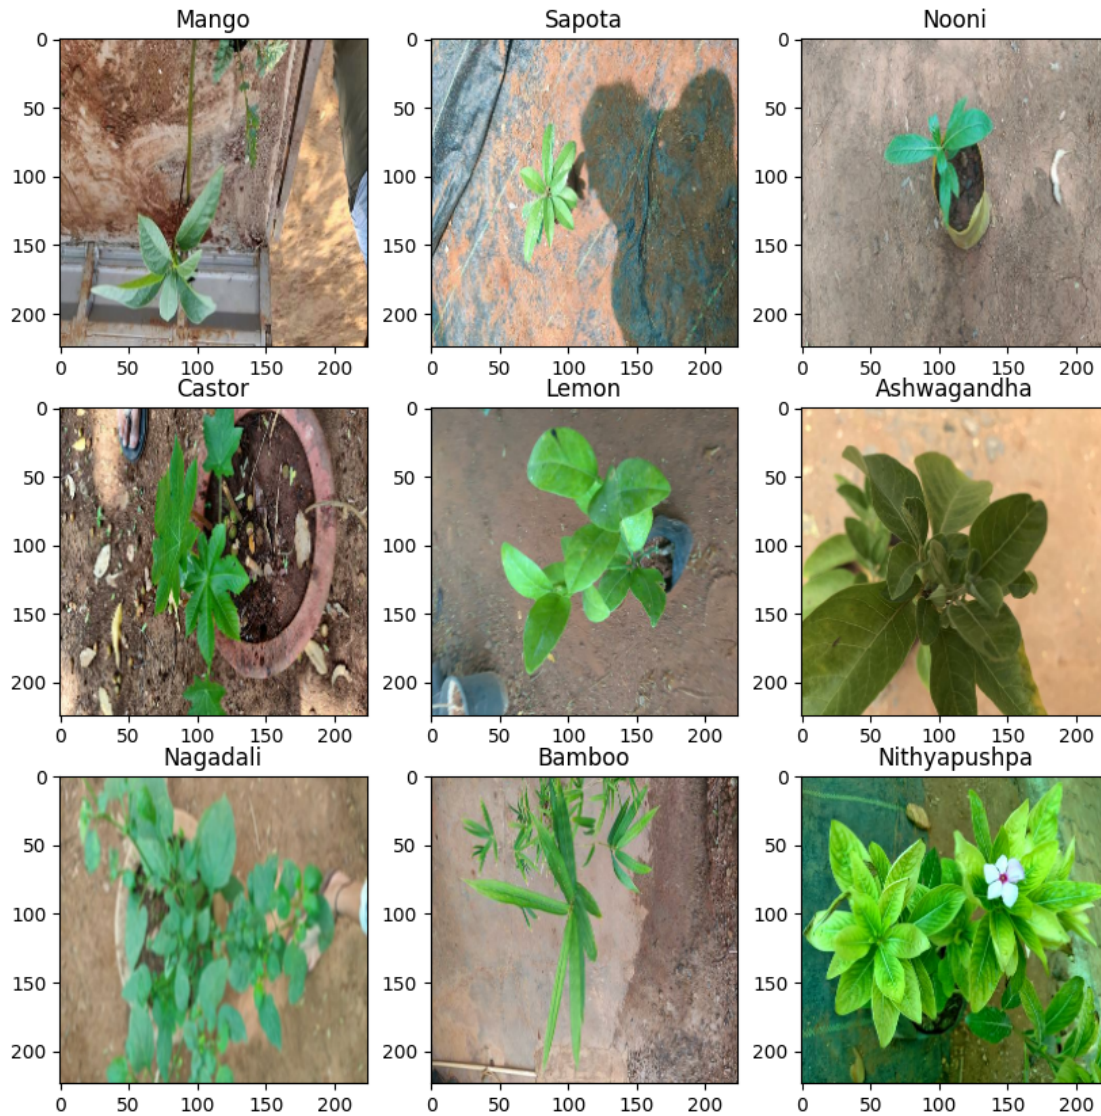
```
[12]: type(class_names)
```

```
[12]: list
```

```
[13]: size = (224,224)
ds_train = ds_train.map(lambda image, label: (tf.image.resize(image,size) ,
↳label))
ds_val = ds_train.map(lambda image, label: (tf.image.resize(image,size) ,
↳label))
```

```
[14]: import matplotlib.pyplot as plt
```

```
[15]: plt.figure(figsize = (10,10))
for images,labels in ds_train.take(1):
    for i in range(9):
        ax = plt.subplot(3,3,i+1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("on")
```



```
[16]: # Initialize an empty list to store images
images_list = []

# Initialize an empty NumPy array for labels
labels_array = np.array([])

# Iterate through the dataset
for image_batch, label_batch in ds_train:
    # Append the image batch to the images list
    images_list.append(image_batch.numpy())

    # Concatenate the label batch to the existing labels_array
    labels_array = np.concatenate([labels_array, label_batch.numpy()], axis=0)
```

```
# Concatenate the list of image batches into a NumPy array
images_array = np.concatenate(images_list, axis=0)

# Now, 'images_array' contains all the images in a 4D NumPy array, and
↪ 'labels_array' contains all the labels as a NumPy array
```

```
[17]: images_array.shape
```

```
[17]: (3567, 224, 224, 3)
```

```
[18]: labels_array
```

```
[18]: array([18., 25., 25., ..., 20.,  1.,  2.])
```

```
[19]: images_array[0]
```

```
[19]: array([[184.7496 , 180.7496 , 177.7496 ],
           [195.97075, 191.97075, 188.97075],
           [199.6789 , 194.6789 , 190.6789 ],
           ...,
           [166.30316, 156.30316, 146.30316],
           [172.80484, 162.80484, 152.80484],
           [169.85579, 159.85579, 149.85579]],

           [[190.51427, 186.20624, 182.89821],
           [196.19077, 191.88274, 188.57469],
           [189.06744, 184.06744, 180.06744],
           ...,
           [167.57985, 157.57985, 147.8879 ],
           [170.82776, 160.82776, 151.13579],
           [170.25853, 160.25853, 150.25853]],

           [[194.62302, 189.62302, 185.62302],
           [194.40865, 189.40865, 185.40865],
           [181.74083, 176.74083, 172.74083],
           ...,
           [168.41255, 158.41255, 149.41255],
           [169.27028, 159.27028, 150.27028],
           [172.21446, 162.21446, 152.21446]],

           ...,

           [[208.2597 , 203.2597 , 199.2597 ],
           [219.00333, 211.00333, 208.00333],
           [214.03223, 206.03223, 203.03223],
           ...,
```

```

[176.33624, 177.51117, 175.89697],
[169.93007, 167.97415, 166.96475],
[164.251 , 162.02573, 159.46089]],

[[216.25323, 209.17728, 205.86926],
[215.01393, 207.01393, 204.01393],
[216.77718, 208.08519, 203.70122],
...,
[168.1254 , 170.1254 , 168.85358],
[180.52542, 178.22838, 179.648 ],
[197.1066 , 193.94708, 189.51361]],

[[225.57068, 217.57068, 214.57068],
[215.28612, 207.28612, 204.28612],
[190.13632, 181.13632, 176.13632],
...,
[173.1664 , 175.1664 , 173.77357],
[172.84341, 170.84341, 172.87906],
[184.34416, 179.54948, 175.15842]]], dtype=float32)

```

```


[20]: plt.imshow(images_array[2].astype("uint8"))
      plt.axis("on")

```

```

[20]: (-0.5, 223.5, 223.5, -0.5)

```



```
[21]: class_names[int(labels_array[2])]
```

```
[21]: 'Lemon_grass'
```

Training the model

```
[22]: from tensorflow.keras.layers import Input,Dense,Flatten
      from tensorflow.keras.models import Model
      from tensorflow.keras.applications.vgg16 import VGG16
      from tensorflow.keras.applications.vgg16 import preprocess_input
      from tensorflow.keras.preprocessing import image
      from tensorflow.keras.preprocessing.image import ImageDataGenerator
      from tensorflow.keras.models import Sequential
```

```
[23]: IMAGE_SIZE = [224,224]
```

```
[24]: vgg = VGG16(input_shape=IMAGE_SIZE+[3] , weights = 'imagenet' , include_top =_
      ↪False)
```

```
[25]: for layer in vgg.layers:
      layer.trainable = False
```

```
[26]: x = Flatten()(vgg.output)
```

```
[27]: prediction = Dense(len(class_names) , activation = "softmax")(x)
```

```
[28]: model = Model(inputs = vgg.input,outputs = prediction)
      model.summary()
```

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584

block2_pool1 (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool1 (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool1 (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool1 (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 40)	1,003,560

Total params: 15,718,248 (59.96 MB)

Trainable params: 1,003,560 (3.83 MB)

Non-trainable params: 14,714,688 (56.13 MB)

```
[29]: model.compile(
    loss = 'categorical_crossentropy',
    optimizer = 'adam',
    metrics = ['accuracy']
)
```

```
[30]: labels_array
```

```
[30]: array([18., 25., 25., ..., 20.,  1.,  2.]
```

```
[31]: # One-hot encode the labels  
labels_one_hot = tf.keras.utils.to_categorical(labels_array, num_classes=40)
```

```
[32]: r = model.fit(  
    images_array,  
    labels_one_hot,  
    batch_size=16,  
    epochs = 5  
)
```

```
Epoch 1/5  
223/223          973s 4s/step -  
accuracy: 0.4400 - loss: 19.5057  
Epoch 2/5  
223/223          1035s 5s/step -  
accuracy: 0.9143 - loss: 2.8841  
Epoch 3/5  
223/223          593s 3s/step -  
accuracy: 0.9565 - loss: 1.3523  
Epoch 4/5  
223/223          455s 2s/step -  
accuracy: 0.9685 - loss: 1.1977  
Epoch 5/5  
223/223          453s 2s/step -  
accuracy: 0.9800 - loss: 0.9391
```

```
[33]: # Save the entire model, including architecture, weights, and optimizer state  
model.save("my_model.h5") # Provide a filename with a '.h5' extension
```


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

Testing model;

```
[34]: import matplotlib.pyplot as plt
```

```
[35]: plt.figure(figsize = (10,10))  
for images,labels in ds_val.take(1):  
    for i in range(9):  
        ax = plt.subplot(3,3,i+1)  
        plt.imshow(images[i].numpy().astype("uint8"))  
        plt.title(class_names[labels[i]])  
        plt.axis("on")
```





```
[36]: import numpy as np
```

```
[37]: # Initialize empty lists to store images and labels
images_list_v = []
labels_list_v = []

# Iterate through the dataset
for image_batch, label_batch in ds_val:
    # Append the image batch to the images list
    images_list_v.append(image_batch.numpy())

    # Append the label batch to the labels list
    labels_list_v.append(label_batch.numpy())
```

```
# Concatenate the list of image batches into a NumPy array
images_array_v = np.concatenate(images_list_v, axis=0)

# Concatenate the list of label batches into a NumPy array
labels_array_v = np.concatenate(labels_list_v, axis=0)

# Now, 'images_array' contains all the images in a 4D NumPy array, and
↳ 'labels_array' contains all the corresponding labels as a NumPy array
```

```
[38]: images_array_v.shape
```

```
[38]: (3567, 224, 224, 3)
```

```
[39]: labels_array_v.shape
```

```
[39]: (3567,)
```

```
[40]: # One-hot encode the labels
labels_one_hot_v = tf.keras.utils.to_categorical(labels_array_v, num_classes=40)
```

```
[41]: model_new = tf.keras.models.load_model("my_model.h5")
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you train or evaluate the model.

```
[42]: model_new.summary()
```

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0

block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 40)	1,003,560

Total params: 15,718,250 (59.96 MB)

Trainable params: 1,003,560 (3.83 MB)

Non-trainable params: 14,714,688 (56.13 MB)

Optimizer params: 2 (12.00 B)

```
[43]: model_new.evaluate(images_array_v, labels_one_hot_v)
```

```
112/112          449s 4s/step -
accuracy: 0.9888 - loss: 0.7650
```

```
[43]: [0.45777466893196106, 0.9873843789100647]
```

```
[44]: y_pred = model_new.predict(images_array_v)
```

112/112                      454s 4s/step

```
[49]: y_pred
```

```
[49]: array([[0., 0., 0., ..., 0., 0., 0.],
          [0., 1., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          ...,
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 1., 0., ..., 0., 0., 0.]], dtype=float32)
```

```
[46]: from sklearn.metrics import confusion_matrix , classification_report
import numpy as np
y_pred_classes = [np.argmax(element) for element in y_pred]

print("Classification Report: \n", classification_report(labels_array_v,
↪y_pred_classes))
```

Classification Report:


	precision	recall	f1-score	support
0	1.00	1.00	1.00	99
1	1.00	1.00	1.00	97
2	0.94	1.00	0.97	94
3	1.00	0.99	0.99	91
4	1.00	1.00	1.00	89
5	0.98	1.00	0.99	86
6	0.94	1.00	0.97	81
7	1.00	1.00	1.00	96
8	1.00	0.98	0.99	87
9	0.97	0.99	0.98	91
10	1.00	0.99	0.99	96
11	1.00	0.99	0.99	83
12	0.96	1.00	0.98	98
13	1.00	1.00	1.00	70
14	0.99	1.00	0.99	90
15	0.98	1.00	0.99	103
16	0.98	0.91	0.95	69
17	0.99	1.00	0.99	88
18	0.99	1.00	0.99	82
19	1.00	0.97	0.98	94
20	1.00	1.00	1.00	88
21	1.00	0.99	0.99	80
22	1.00	0.89	0.94	87
23	0.99	0.99	0.99	111

24	0.99	1.00	0.99	85
25	1.00	1.00	1.00	89
26	0.96	1.00	0.98	89
27	1.00	1.00	1.00	88
28	0.95	0.98	0.97	86
29	1.00	0.99	0.99	93
30	1.00	1.00	1.00	88
31	1.00	1.00	1.00	89
32	0.97	1.00	0.98	86
33	0.95	1.00	0.97	96
34	1.00	0.97	0.98	89
35	1.00	0.94	0.97	85
36	1.00	0.99	0.99	97
37	1.00	1.00	1.00	85
38	1.00	0.98	0.99	85
39	0.99	0.95	0.97	87
accuracy				0.99 3567
macro avg				0.99 0.99 0.99 3567
weighted avg				0.99 0.99 0.99 3567

```
[56]: import seaborn as sns

plt.figure(figsize=(15,15))
sns.heatmap(confusion_matrix(labels_array_v, y_pred_classes))
```

```
[56]: <Axes: >
```



[ ]: